

# Online Retail Data Clustering

February 19, 2026

## 1 IMPORTING THE REQUIRED LIBRARIES

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from sklearn.preprocessing import StandardScaler

# Setting to make numbers easier to read on display
pd.options.display.float_format = '{:20.2f}'.format

# Show all columns on output
pd.set_option('display.max_columns', 999)
```

```
[2]: df = pd.read_excel("online_retail_II.xlsx", sheet_name=0)

df.head(10)
```

```
[2]:   Invoice StockCode          Description  Quantity \
0    489434    85048  15CM CHRISTMAS GLASS BALL 20 LIGHTS      12
1    489434    79323P           PINK CHERRY LIGHTS      12
2    489434    79323W           WHITE CHERRY LIGHTS      12
3    489434    22041        RECORD FRAME 7" SINGLE SIZE      48
4    489434    21232  STRAWBERRY CERAMIC TRINKET BOX      24
5    489434    22064           PINK DOUGHNUT TRINKET POT      24
6    489434    21871           SAVE THE PLANET MUG      24
7    489434    21523  FANCY FONT HOME SWEET HOME DOORMAT      10
8    489435    22350                 CAT BOWL      12
9    489435    22349  DOG BOWL , CHASING BALL DESIGN      12

          InvoiceDate      Price  Customer ID \
0  2009-12-01 07:45:00    6.95     13085.00
1  2009-12-01 07:45:00    6.75     13085.00
2  2009-12-01 07:45:00    6.75     13085.00
3  2009-12-01 07:45:00    2.10     13085.00
```

```
4 2009-12-01 07:45:00           1.25      13085.00
5 2009-12-01 07:45:00           1.65      13085.00
6 2009-12-01 07:45:00           1.25      13085.00
7 2009-12-01 07:45:00           5.95      13085.00
8 2009-12-01 07:46:00           2.55      13085.00
9 2009-12-01 07:46:00           3.75      13085.00
```

```
          Country
0  United Kingdom
1  United Kingdom
2  United Kingdom
3  United Kingdom
4  United Kingdom
5  United Kingdom
6  United Kingdom
7  United Kingdom
8  United Kingdom
9  United Kingdom
```

## 2 EXPLORATORY DATA ANALYSIS

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 525461 entries, 0 to 525460
Data columns (total 8 columns):
 #   Column        Non-Null Count  Dtype  
--- 
 0   Invoice       525461 non-null   object 
 1   StockCode     525461 non-null   object 
 2   Description   522533 non-null   object 
 3   Quantity      525461 non-null   int64  
 4   InvoiceDate   525461 non-null   datetime64[ns]
 5   Price         525461 non-null   float64
 6   Customer ID  417534 non-null   float64
 7   Country       525461 non-null   object 
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 32.1+ MB
```

```
[4]: df.describe()
```

```
          Quantity           InvoiceDate \
count      525461.00                  525461
mean       10.34  2010-06-28 11:37:36.845017856
min      -9600.00                  2009-12-01 07:45:00
25%        1.00                  2010-03-21 12:20:00
50%        3.00                  2010-07-06 09:51:00
```

```

75%           10.00      2010-10-15 12:45:00
max          19152.00    2010-12-09 20:01:00
std           107.42      NaN

```

	Price	Customer ID
count	525461.00	417534.00
mean	4.69	15360.65
min	-53594.36	12346.00
25%	1.25	13983.00
50%	2.10	15311.00
75%	4.21	16799.00
max	25111.09	18287.00
std	146.13	1680.81

[9]: df.describe(include="O")

	Invoice	StockCode	Description	Country
count	525461	525461	522533	525461
unique	28816	4632	4681	40
top	537434	85123A	WHITE HANGING HEART T-LIGHT HOLDER	United Kingdom
freq	675	3516	3549	485852

[10]: df[df["Customer ID"].isna()].head(10)

	Invoice	StockCode	Description	Quantity
263	489464	21733	85123a mixed	-96
283	489463	71477	short	-240
284	489467	85123A	21733 mixed	-192
470	489521	21646	Nan	-50
577	489525	85226C	BLUE PULL BACK RACING CAR	1
578	489525	85227	SET/6 3D KIT CARDS FOR KIDS	1
1055	489548	22271	FELTCRAFT DOLL ROSIE	1
1056	489548	22254	FELT TOADSTOOL LARGE	12
1057	489548	22273	FELTCRAFT DOLL MOLLY	3
1058	489548	22195	LARGE HEART MEASURING SPOONS	1

	InvoiceDate	Price	Customer ID	Country
263	2009-12-01 10:52:00	0.00	NaN	United Kingdom
283	2009-12-01 10:52:00	0.00	NaN	United Kingdom
284	2009-12-01 10:53:00	0.00	NaN	United Kingdom
470	2009-12-01 11:44:00	0.00	NaN	United Kingdom
577	2009-12-01 11:49:00	0.55	NaN	United Kingdom
578	2009-12-01 11:49:00	0.85	NaN	United Kingdom
1055	2009-12-01 12:32:00	2.95	NaN	United Kingdom
1056	2009-12-01 12:32:00	1.25	NaN	United Kingdom
1057	2009-12-01 12:32:00	2.95	NaN	United Kingdom
1058	2009-12-01 12:32:00	1.65	NaN	United Kingdom

```
[11]: df[df["Quantity"]<0].head(10)
```

```
[11]:      Invoice StockCode          Description  Quantity \
178    C489449    22087    PAPER BUNTING WHITE LACE     -12
179    C489449    85206A   CREAM FELT EASTER EGG BASKET     -6
180    C489449    21895  POTTING SHED SOW 'N' GROW SET     -4
181    C489449    21896    POTTING SHED TWINE     -6
182    C489449    22083  PAPER CHAIN KIT RETRO SPOT     -12
183    C489449    21871        SAVE THE PLANET MUG     -12
184    C489449    84946  ANTIQUE SILVER TEA GLASS ETCHED     -12
185    C489449    84970S HANGING HEART ZINC T-LIGHT HOLDER    -24
186    C489449    22090  PAPER BUNTING RETRO SPOTS     -12
196    C489459    90200A   PURPLE SWEETHEART BRACELET     -3

      InvoiceDate          Price  Customer ID \
178 2009-12-01 10:33:00      2.95    16321.00
179 2009-12-01 10:33:00      1.65    16321.00
180 2009-12-01 10:33:00      4.25    16321.00
181 2009-12-01 10:33:00      2.10    16321.00
182 2009-12-01 10:33:00      2.95    16321.00
183 2009-12-01 10:33:00      1.25    16321.00
184 2009-12-01 10:33:00      1.25    16321.00
185 2009-12-01 10:33:00      0.85    16321.00
186 2009-12-01 10:33:00      2.95    16321.00
196 2009-12-01 10:44:00      4.25    17592.00

      Country
178    Australia
179    Australia
180    Australia
181    Australia
182    Australia
183    Australia
184    Australia
185    Australia
186    Australia
196  United Kingdom
```

```
[13]: df["Invoice"] = df["Invoice"].astype("str")
df[df["Invoice"].str.match("\d{6}$") == False]
```

```
[13]:      Invoice StockCode          Description  Quantity \
178    C489449    22087    PAPER BUNTING WHITE LACE     -12
179    C489449    85206A   CREAM FELT EASTER EGG BASKET     -6
180    C489449    21895  POTTING SHED SOW 'N' GROW SET     -4
181    C489449    21896    POTTING SHED TWINE     -6
182    C489449    22083  PAPER CHAIN KIT RETRO SPOT     -12
```

```

...
524695 C538123    22956      36 FOIL HEART CAKE CASES     -2
524696 C538124        M          Manual                 -4
524697 C538124    22699 ROSES REGENCY TEACUP AND SAUCER     -1
524698 C538124    22423      REGENCY CAKESTAND 3 TIER     -1
525282 C538164    35004B      SET OF 3 BLACK FLYING DUCKS     -1

                    InvoiceDate          Price      Customer ID \
178    2009-12-01 10:33:00      2.95      16321.00
179    2009-12-01 10:33:00      1.65      16321.00
180    2009-12-01 10:33:00      4.25      16321.00
181    2009-12-01 10:33:00      2.10      16321.00
182    2009-12-01 10:33:00      2.95      16321.00
...
524695 2010-12-09 15:41:00      2.10      12605.00
524696 2010-12-09 15:43:00      0.50      15329.00
524697 2010-12-09 15:43:00      2.95      15329.00
524698 2010-12-09 15:43:00     12.75      15329.00
525282 2010-12-09 17:32:00      1.95      14031.00

                    Country
178      Australia
179      Australia
180      Australia
181      Australia
182      Australia
...
524695      Germany
524696  United Kingdom
524697  United Kingdom
524698  United Kingdom
525282  United Kingdom

```

[10209 rows x 8 columns]

[14]: df["Invoice"].str.replace("[0-9]", "", regex=True).unique()

[14]: array(['', 'C', 'A'], dtype=object)

[16]: df[df["Invoice"].str.startswith("A")]

```

[16]:      Invoice StockCode      Description  Quantity      InvoiceDate \
179403   A506401      B  Adjust bad debt      1 2010-04-29 13:36:00
276274   A516228      B  Adjust bad debt      1 2010-07-19 11:24:00
403472   A528059      B  Adjust bad debt      1 2010-10-20 12:04:00

```

Price Customer ID Country

```

179403          -53594.36      NaN  United Kingdom
276274          -44031.79      NaN  United Kingdom
403472          -38925.87      NaN  United Kingdom

```

```
[21]: df["StockCode"] = df["StockCode"].astype("str")
df[(df["StockCode"].str.match("^\\d{5}$") == False) & (df["StockCode"].str.
    match("^\\d{5}[a-zA-Z]+$") == False)]["StockCode"].unique()
```

```
[21]: array(['POST', 'D', 'DCGS0058', 'DCGS0068', 'DOT', 'M', 'DCGS0004',
       'DCGS0076', 'C2', 'BANK CHARGES', 'DCGS0003', 'TEST001',
       'gift_0001_80', 'DCGS0072', 'gift_0001_20', 'DCGS0044', 'TEST002',
       'gift_0001_10', 'gift_0001_50', 'DCGS0066N', 'gift_0001_30',
       'PADS', 'ADJUST', 'gift_0001_40', 'gift_0001_60', 'gift_0001_70',
       'gift_0001_90', 'DCGSSGIRL', 'DCGS0006', 'DCGS0016', 'DCGS0027',
       'DCGS0036', 'DCGS0039', 'DCGS0060', 'DCGS0056', 'DCGS0059', 'GIFT',
       'DCGSLBOY', 'm', 'DCGS0053', 'DCGS0062', 'DCGS0037', 'DCGSSBOY',
       'DCGSLGIRL', 'S', 'DCGS0069', 'DCGS0070', 'DCGS0075', 'B',
       'DCGS0041', 'ADJUST2', '47503J ', 'C3', 'SP1002', 'AMAZONFEE'],
      dtype=object)
```

```
[22]: df[df["StockCode"].str.contains("^\u00c4DOT")]
```

	Invoice	StockCode	Description	Quantity	InvoiceDate
2379	489597	DOT	DOTCOM POSTAGE	1	2009-12-01 14:28:00
2539	489600	DOT	DOTCOM POSTAGE	1	2009-12-01 14:43:00
2551	489601	DOT	DOTCOM POSTAGE	1	2009-12-01 14:44:00
2571	489602	DOT	DOTCOM POSTAGE	1	2009-12-01 14:45:00
2619	489603	DOT	DOTCOM POSTAGE	1	2009-12-01 14:46:00
...	...	...	...	...	...
524272	538071	DOT	DOTCOM POSTAGE	1	2010-12-09 14:09:00
524887	538148	DOT	DOTCOM POSTAGE	1	2010-12-09 16:26:00
525000	538149	DOT	DOTCOM POSTAGE	1	2010-12-09 16:27:00
525126	538153	DOT	DOTCOM POSTAGE	1	2010-12-09 16:31:00
525147	538154	DOT	DOTCOM POSTAGE	1	2010-12-09 16:35:00
		Price	Customer ID	Country	
2379		647.19	NaN	United Kingdom	
2539		55.96	NaN	United Kingdom	
2551		68.39	NaN	United Kingdom	
2571		59.35	NaN	United Kingdom	
2619		42.39	NaN	United Kingdom	
...	...	...	...	...	
524272		885.94	NaN	United Kingdom	
524887		547.32	NaN	United Kingdom	
525000		620.68	NaN	United Kingdom	
525126		822.94	NaN	United Kingdom	
525147		85.79	NaN	United Kingdom	

[736 rows x 8 columns]

### 3 DATA CLEANING

```
[23]: cleaned_df = df.copy()
```

```
[24]: cleaned_df["Invoice"] = cleaned_df["Invoice"].astype("str")
mask = (
    cleaned_df["Invoice"].str.match("^\\d{6}$") == True
)
cleaned_df = cleaned_df[mask]
cleaned_df
```

```
[24]:      Invoice StockCode          Description  Quantity \
0        489434    85048  15CM CHRISTMAS GLASS BALL 20 LIGHTS     12
1        489434    79323P           PINK CHERRY LIGHTS     12
2        489434    79323W           WHITE CHERRY LIGHTS     12
3        489434    22041      RECORD FRAME 7" SINGLE SIZE     48
4        489434    21232  STRAWBERRY CERAMIC TRINKET BOX     24
...
525456   538171    22271      FELTCRAFT DOLL ROSIE      2
525457   538171    22750      FELTCRAFT PRINCESS LOLA DOLL      1
525458   538171    22751      FELTCRAFT PRINCESS OLIVIA DOLL      1
525459   538171    20970  PINK FLORAL FELTCRAFT SHOULDER BAG      2
525460   538171    21931      JUMBO STORAGE BAG SUKI      2
```

```
      InvoiceDate      Price  Customer ID \
0  2009-12-01 07:45:00    6.95    13085.00
1  2009-12-01 07:45:00    6.75    13085.00
2  2009-12-01 07:45:00    6.75    13085.00
3  2009-12-01 07:45:00    2.10    13085.00
4  2009-12-01 07:45:00    1.25    13085.00
...
525456 2010-12-09 20:01:00    2.95    17530.00
525457 2010-12-09 20:01:00    3.75    17530.00
525458 2010-12-09 20:01:00    3.75    17530.00
525459 2010-12-09 20:01:00    3.75    17530.00
525460 2010-12-09 20:01:00    1.95    17530.00
```

```
      Country
0  United Kingdom
1  United Kingdom
2  United Kingdom
3  United Kingdom
4  United Kingdom
...
...
```

```
525456 United Kingdom  
525457 United Kingdom  
525458 United Kingdom  
525459 United Kingdom  
525460 United Kingdom
```

[515252 rows x 8 columns]

```
[26]: cleaned_df["StockCode"] = cleaned_df["StockCode"].astype("str")  
mask = (  
    (cleaned_df["StockCode"].str.match("^\\d{5}$") == True)  
    | (cleaned_df["StockCode"].str.match("^\\d{5}[a-zA-Z]+$") == True)  
    | (cleaned_df["StockCode"].str.match("^PADS$") == True)  
)  
  
cleaned_df = cleaned_df[mask]  
cleaned_df
```

C:\Users\udayd\AppData\Local\Temp\ipykernel\_8468\3378931196.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
cleaned_df["StockCode"] = cleaned_df["StockCode"].astype("str")
```

```
[26]:  
      Invoice StockCode          Description  Quantity  \\\n0       489434     85048  15CM CHRISTMAS GLASS BALL 20 LIGHTS      12  
1       489434     79323P           PINK CHERRY LIGHTS      12  
2       489434     79323W           WHITE CHERRY LIGHTS      12  
3       489434     22041        RECORD FRAME 7" SINGLE SIZE      48  
4       489434     21232   STRAWBERRY CERAMIC TRINKET BOX      24  
...  
525456  538171     22271           FELTCRAFT DOLL ROSIE      2  
525457  538171     22750           FELTCRAFT PRINCESS LOLA DOLL      1  
525458  538171     22751           FELTCRAFT PRINCESS OLIVIA DOLL      1  
525459  538171     20970  PINK FLORAL FELTCRAFT SHOULDER BAG      2  
525460  538171     21931           JUMBO STORAGE BAG SUKI      2  
  
      InvoiceDate          Price  Customer ID  \\\n0  2009-12-01 07:45:00      6.95  13085.00  
1  2009-12-01 07:45:00      6.75  13085.00  
2  2009-12-01 07:45:00      6.75  13085.00  
3  2009-12-01 07:45:00      2.10  13085.00  
4  2009-12-01 07:45:00      1.25  13085.00  
...  
525456  2010-12-09 20:01:00      2.95  17530.00
```

```

525457 2010-12-09 20:01:00           3.75      17530.00
525458 2010-12-09 20:01:00           3.75      17530.00
525459 2010-12-09 20:01:00           3.75      17530.00
525460 2010-12-09 20:01:00           1.95      17530.00

          Country
0    United Kingdom
1    United Kingdom
2    United Kingdom
3    United Kingdom
4    United Kingdom
...
525456  United Kingdom
525457  United Kingdom
525458  United Kingdom
525459  United Kingdom
525460  United Kingdom

```

[512796 rows x 8 columns]

[28]: `cleaned_df.dropna(subset = ["Customer ID"], inplace = True)`

```
C:\Users\udayd\AppData\Local\Temp\ipykernel_8468\2119709935.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
`cleaned_df.dropna(subset = ["Customer ID"], inplace = True)`

[29]: `cleaned_df.describe()`

	Quantity	InvoiceDate	\
count	406337.00	406337	
mean	13.62	2010-07-01 10:11:06.543288320	
min	1.00	2009-12-01 07:45:00	
25%	2.00	2010-03-26 14:01:00	
50%	5.00	2010-07-09 15:48:00	
75%	12.00	2010-10-14 17:09:00	
max	19152.00	2010-12-09 20:01:00	
std	97.00	NaN	

	Price	Customer ID
count	406337.00	406337.00
mean	2.99	15373.63
min	0.00	12346.00
25%	1.25	14004.00
50%	1.95	15326.00

```
75%           3.75          16814.00
max           295.00         18287.00
std            4.29          1677.37
```

```
[31]: len(cleaned_df[cleaned_df["Price"] == 0])
```

```
[31]: 28
```

```
[34]: cleaned_df = cleaned_df[cleaned_df["Price"] > 0.0]
```

```
[35]: cleaned_df.describe()
```

```
[35]:
```

	Quantity	InvoiceDate
count	406309.00	406309
mean	13.62	2010-07-01 10:14:25.869572352
min	1.00	2009-12-01 07:45:00
25%	2.00	2010-03-26 14:01:00
50%	5.00	2010-07-09 15:48:00
75%	12.00	2010-10-14 17:09:00
max	19152.00	2010-12-09 20:01:00
std	97.00	NaN

	Price	Customer ID
count	406309.00	406309.00
mean	2.99	15373.72
min	0.00	12346.00
25%	1.25	14006.00
50%	1.95	15326.00
75%	3.75	16814.00
max	295.00	18287.00
std	4.29	1677.33

```
[36]: cleaned_df["Price"].min()
```

```
[36]: 0.001
```

```
[37]: len(cleaned_df)/len(df)
```

```
[37]: 0.7732429238325965
```

Dropped about 23% of records.

## 4 FEATURE ENGINEERING

```
[38]: cleaned_df["SalesLineTotal"] = cleaned_df["Quantity"] * cleaned_df["Price"]

cleaned_df
```

```
[38]:      Invoice StockCode          Description  Quantity \
0        489434    85048  15CM CHRISTMAS GLASS BALL 20 LIGHTS      12
1        489434    79323P           PINK CHERRY LIGHTS      12
2        489434    79323W           WHITE CHERRY LIGHTS      12
3        489434    22041  RECORD FRAME 7" SINGLE SIZE       48
4        489434    21232  STRAWBERRY CERAMIC TRINKET BOX      24
...
525456  538171    22271  FELTCRAFT DOLL ROSIE          ...
525457  538171    22750  FELTCRAFT PRINCESS LOLA DOLL      1
525458  538171    22751  FELTCRAFT PRINCESS OLIVIA DOLL      1
525459  538171    20970  PINK FLORAL FELTCRAFT SHOULDER BAG      2
525460  538171    21931  JUMBO STORAGE BAG SUKI          ...

      InvoiceDate      Price  Customer ID \
0  2009-12-01 07:45:00    6.95  13085.00
1  2009-12-01 07:45:00    6.75  13085.00
2  2009-12-01 07:45:00    6.75  13085.00
3  2009-12-01 07:45:00    2.10  13085.00
4  2009-12-01 07:45:00    1.25  13085.00
...
525456 2010-12-09 20:01:00    ...
525457 2010-12-09 20:01:00    2.95  17530.00
525458 2010-12-09 20:01:00    3.75  17530.00
525459 2010-12-09 20:01:00    3.75  17530.00
525460 2010-12-09 20:01:00    1.95  17530.00

      Country  SalesLineTotal
0  United Kingdom      83.40
1  United Kingdom      81.00
2  United Kingdom      81.00
3  United Kingdom     100.80
4  United Kingdom      30.00
...
525456  United Kingdom      ...
525457  United Kingdom      5.90
525458  United Kingdom      3.75
525459  United Kingdom      3.75
525460  United Kingdom      7.50
525460  United Kingdom      3.90
```

[406309 rows x 9 columns]

```
[39]: aggregated_df = cleaned_df.groupby(by="Customer ID", as_index=False) \
    .agg(
        MonetaryValue=("SalesLineTotal", "sum"),
        Frequency=("Invoice", "nunique"),
        LastInvoiceDate=("InvoiceDate", "max")
    )
```

```
aggregated_df.head(5)
```

```
[39]:      Customer ID      MonetaryValue  Frequency      LastInvoiceDate
0          12346.00           169.36        2 2010-06-28 13:53:00
1          12347.00           1323.32       2 2010-12-07 14:57:00
2          12348.00            221.16       1 2010-09-27 14:59:00
3          12349.00           2221.14       2 2010-10-28 08:23:00
4          12351.00           300.93        1 2010-11-29 15:23:00
```

```
[40]: max_invoice_date = aggregated_df["LastInvoiceDate"].max()
```

```
aggregated_df["Recency"] = (max_invoice_date -  
                             aggregated_df["LastInvoiceDate"]).dt.days
```

```
aggregated_df.head(5)
```

```
[40]:      Customer ID      MonetaryValue  Frequency      LastInvoiceDate \
0          12346.00           169.36        2 2010-06-28 13:53:00
1          12347.00           1323.32       2 2010-12-07 14:57:00
2          12348.00            221.16       1 2010-09-27 14:59:00
3          12349.00           2221.14       2 2010-10-28 08:23:00
4          12351.00           300.93        1 2010-11-29 15:23:00

      Recency
0         164
1          2
2         73
3         42
4         10
```

```
[41]: plt.figure(figsize=(15, 5))
```

```
plt.subplot(1, 3, 1)
plt.hist(aggregated_df['MonetaryValue'], bins=10, color='skyblue',  

         edgecolor='black')
plt.title('Monetary Value Distribution')
plt.xlabel('Monetary Value')
plt.ylabel('Count')

plt.subplot(1, 3, 2)
plt.hist(aggregated_df['Frequency'], bins=10, color='lightgreen',  


         edgecolor='black')
plt.title('Frequency Distribution')
plt.xlabel('Frequency')
plt.ylabel('Count')
```

```

plt.subplot(1, 3, 1)
plt.hist(aggregated_df['Recency'], bins=20, color='salmon', edgecolor='black')
plt.title('Recency Distribution')
plt.xlabel('Recency')
plt.ylabel('Count')

plt.tight_layout()
plt.show()

```



```

[42]: plt.figure(figsize=(15, 5))


plt.subplot(1, 3, 1)
sns.boxplot(data=aggregated_df['MonetaryValue'], color='skyblue')
plt.title('Monetary Value Boxplot')
plt.xlabel('Monetary Value')

plt.subplot(1, 3, 2)
sns.boxplot(data=aggregated_df['Frequency'], color='lightgreen')
plt.title('Frequency Boxplot')
plt.xlabel('Frequency')

plt.subplot(1, 3, 3)
sns.boxplot(data=aggregated_df['Recency'], color='salmon')
plt.title('Recency Boxplot')
plt.xlabel('Recency')

plt.tight_layout()
plt.show()

```



```
[43]: M_Q1 = aggregated_df[["MonetaryValue"]].quantile(0.25)
M_Q3 = aggregated_df[["MonetaryValue"]].quantile(0.75)
M_IQR = M_Q3 - M_Q1

monetary_outliers_df = aggregated_df[(aggregated_df[["MonetaryValue"]] > (M_Q3 + 1.5 * M_IQR)) | (aggregated_df[["MonetaryValue"]] < (M_Q1 - 1.5 * M_IQR))].
copy()

monetary_outliers_df.describe()
```

	Customer ID	MonetaryValue	Frequency	\
count	423.00	423.00	423.00	
mean	15103.04	12188.10	17.17	
min	12357.00	3802.04	1.00	
25%	13622.00	4605.94	8.00	
50%	14961.00	6191.32	12.00	
75%	16692.00	10273.24	18.00	
max	18260.00	349164.35	183.00	
std	1728.66	25830.85	19.73	
	LastInvoiceDate		Recency	
count	423		423.00	
mean	2010-11-09 12:26:02.978723328		30.04	
min	2009-12-10 18:03:00		0.00	
25%	2010-11-08 13:17:30		3.00	
50%	2010-11-26 12:19:00		13.00	
75%	2010-12-06 10:34:30		31.00	
max	2010-12-09 19:32:00		364.00	
std	NaN		51.54	

```
[44]: F_Q1 = aggregated_df['Frequency'].quantile(0.25)
F_Q3 = aggregated_df['Frequency'].quantile(0.75)
F_IQR = F_Q3 - F_Q1
```

```

frequency_outliers_df = aggregated_df[(aggregated_df['Frequency'] > (F_Q3 + 1.5 * F_IQR)) | (aggregated_df['Frequency'] < (F_Q1 - 1.5 * F_IQR))].copy()

frequency_outliers_df.describe()

```

[44]:

	Customer ID	MonetaryValue	Frequency
count	279.00	279.00	279.00
mean	15352.66	14409.71	23.81
min	12437.00	1094.39	12.00
25%	13800.00	4331.56	13.00
50%	15465.00	6615.77	17.00
75%	16828.50	11692.41	23.00
max	18260.00	349164.35	183.00
std	1748.43	31381.74	21.93

	LastInvoiceDate	Recency
count	279	279.00
mean	2010-11-23 11:06:20.645161216	16.09
min	2010-05-12 16:51:00	0.00
25%	2010-11-20 13:14:30	2.00
50%	2010-12-02 10:46:00	7.00
75%	2010-12-07 11:08:30	19.00
max	2010-12-09 19:32:00	211.00
std	NaN	26.59

[47]:

```

non_outliers_df = aggregated_df[(~aggregated_df.index.isin(monetary_outliers_df.index)) & (~aggregated_df.index.isin(frequency_outliers_df.index))]

```

non\_outliers\_df.describe()

[47]:

	Customer ID	MonetaryValue	Frequency
count	3809.00	3809.00	3809.00
mean	15376.48	885.50	2.86
min	12346.00	1.55	1.00
25%	13912.00	279.91	1.00
50%	15389.00	588.05	2.00
75%	16854.00	1269.05	4.00
max	18287.00	3788.21	11.00
std	1693.20	817.67	2.24

	LastInvoiceDate	Recency
count	3809	3809.00
mean	2010-09-03 11:16:46.516146176	97.08
min	2009-12-01 10:49:00	0.00
25%	2010-07-08 14:48:00	22.00
50%	2010-10-12 16:25:00	58.00
75%	2010-11-17 13:14:00	154.00

```

max           2010-12-09 20:01:00
std                  NaN

```

```

373.00
98.11

```

```

[48]: plt.figure(figsize=(15, 5))


plt.subplot(1, 3, 1)
sns.boxplot(data=non_outliers_df['MonetaryValue'], color='skyblue')
plt.title('Monetary Value Boxplot')
plt.xlabel('Monetary Value')

plt.subplot(1, 3, 2)
sns.boxplot(data=non_outliers_df['Frequency'], color='lightgreen')
plt.title('Frequency Boxplot')
plt.xlabel('Frequency')

plt.subplot(1, 3, 3)
sns.boxplot(data=non_outliers_df['Recency'], color='salmon')
plt.title('Recency Boxplot')
plt.xlabel('Recency')

plt.tight_layout()
plt.show()

```



```

[49]: fig = plt.figure(figsize=(8, 8))

ax = fig.add_subplot(projection="3d")

scatter = ax.scatter(non_outliers_df["MonetaryValue"], non_outliers_df["Frequency"], non_outliers_df["Recency"])


ax.set_xlabel('Monetary Value')
ax.set_ylabel('Frequency')
ax.set_zlabel('Recency')

```

```
ax.set_title('3D Scatter Plot of Customer Data')

plt.show()
```

3D Scatter Plot of Customer Data



```
[50]: scaler = StandardScaler()

scaled_data = scaler.fit_transform(non_outliers_df[["MonetaryValue", "Frequency", "Recency"]])

scaled_data
```

```
[50]: array([[-0.87594534, -0.38488934,  0.68214853],
   [ 0.5355144 , -0.38488934, -0.96925093],
   [-0.81258645, -0.83063076, -0.24548944],
   ...,
   [-0.62197163, -0.83063076,  2.01753946],
   [ 0.44146683, -0.38488934,  0.14187587],
   [ 1.72488781,  0.50659348, -0.81634357]])
```

```
[51]: scaled_data_df = pd.DataFrame(scaled_data, index=non_outliers_df.index, columns=["MonetaryValue", "Frequency", "Recency"))

scaled_data_df
```

```
[51]:      MonetaryValue        Frequency       Recency
0            -0.88           -0.38          0.68
1             0.54           -0.38         -0.97
2            -0.81           -0.83         -0.25
3             1.63           -0.38         -0.56
4            -0.72           -0.83         -0.89
...
4280          -0.30            1.40         -0.82
4281          -0.58           -0.83         -0.32
4282          -0.62           -0.83          2.02
4283            0.44           -0.38          0.14
4284            1.72            0.51         -0.82

[3809 rows x 3 columns]
```

```
[52]: fig = plt.figure(figsize=(8, 8))

ax = fig.add_subplot(projection="3d")


scatter = ax.scatter(scaled_data_df["MonetaryValue"], scaled_data_df["Frequency"], scaled_data_df["Recency"])

ax.set_xlabel('Monetary Value')
ax.set_ylabel('Frequency')
ax.set_zlabel('Recency')

ax.set_title('3D Scatter Plot of Customer Data')

plt.show()
```

3D Scatter Plot of Customer Data



## 5 KMEANS CLUSTERING

```
[53]: max_k = 12

inertia = []
silhouette_scores = []
k_values = range(2, max_k + 1)

for k in k_values:

    kmeans = KMeans(n_clusters=k, random_state=42, max_iter=1000)
```

```

cluster_labels = kmeans.fit_predict(scaled_data_df)

sil_score = silhouette_score(scaled_data_df, cluster_labels)

silhouette_scores.append(sil_score)

inertia.append(kmeans.inertia_)


plt.figure(figsize=(14, 6))

plt.subplot(1, 2, 1)
plt.plot(k_values, inertia, marker='o')
plt.title('KMeans Inertia for Different Values of k')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Inertia')
plt.xticks(k_values)
plt.grid(True)

plt.subplot(1, 2, 2)
plt.plot(k_values, silhouette_scores, marker='o', color='orange')
plt.title('Silhouette Scores for Different Values of k')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Silhouette Score')
plt.xticks(k_values)
plt.grid(True)

plt.tight_layout()
plt.show()

```



```
[54]: kmeans = KMeans(n_clusters=4, random_state=42, max_iter=1000)

cluster_labels = kmeans.fit_predict(scaled_data_df)

cluster_labels
```

[54]: array([1, 0, 2, ..., 1, 0, 0], dtype=int32)

```
[55]: non_outliers_df["Cluster"] = cluster_labels

non_outliers_df
```

C:\Users\udayd\AppData\Local\Temp\ipykernel\_8468\3577770544.py:1:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead  
See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
non\_outliers\_df["Cluster"] = cluster\_labels

```
[55]:      Customer ID    MonetaryValue  Frequency  LastInvoiceDate \
0          12346.00        169.36       2 2010-06-28 13:53:00
1          12347.00        1323.32      2 2010-12-07 14:57:00
2          12348.00        221.16       1 2010-09-27 14:59:00
3          12349.00        2221.14      2 2010-10-28 08:23:00
4          12351.00        300.93       1 2010-11-29 15:23:00
...          ...
4280         18283.00        641.77       6 2010-11-22 15:30:00
4281         18284.00        411.68       1 2010-10-04 11:33:00
4282         18285.00        377.00       1 2010-02-17 10:24:00
4283         18286.00        1246.43      2 2010-08-20 11:57:00
4284         18287.00        2295.71      4 2010-11-22 11:51:00

      Recency  Cluster
0          164       1
1           2       0
2           73       2
3           42       0
4           10       2
...          ...
4280         17       0
4281         66       2
4282         295      1
4283         111      0
4284         17       0

[3809 rows x 6 columns]
```

```
[56]: cluster_colors = {0: '#1f77b4', # Blue
                      1: '#ff7f0e', # Orange
                      2: '#2ca02c', # Green
                      3: '#d62728'} # Red

colors = non_outliers_df['Cluster'].map(cluster_colors)

fig = plt.figure(figsize=(10, 10))
ax = fig.add_subplot(projection='3d')


scatter = ax.scatter(non_outliers_df['MonetaryValue'],
                     non_outliers_df['Frequency'],
                     non_outliers_df['Recency'],
                     c=colors, # Use mapped solid colors
                     marker='o')

ax.set_xlabel('Monetary Value')
ax.set_ylabel('Frequency')
ax.set_zlabel('Recency')

ax.set_title('3D Scatter Plot of Customer Data by Cluster')

plt.show()
```

3D Scatter Plot of Customer Data by Cluster




```
[57]: plt.figure(figsize=(12, 18))

plt.subplot(3, 1, 1)
sns.violinplot(x=non_outliers_df['Cluster'], y=non_outliers_df['MonetaryValue'], palette=cluster_colors, hue=non_outliers_df["Cluster"])
sns.violinplot(y=non_outliers_df['MonetaryValue'], color='gray', linewidth=1.0)
plt.title('Monetary Value by Cluster')
plt.ylabel('Monetary Value')

plt.subplot(3, 1, 2)
```

```
sns.violinplot(x=non_outliers_df['Cluster'], y=non_outliers_df['Frequency'],  
    ↪palette=cluster_colors, hue=non_outliers_df["Cluster"])  
sns.violinplot(y=non_outliers_df['Frequency'], color='gray', linewidth=1.0)  
plt.title('Frequency by Cluster')  
plt.ylabel('Frequency')  
  
plt.subplot(3, 1, 3)  
sns.violinplot(x=non_outliers_df['Cluster'], y=non_outliers_df['Recency'],  
    ↪palette=cluster_colors, hue=non_outliers_df["Cluster"])  
sns.violinplot(y=non_outliers_df['Recency'], color='gray', linewidth=1.0)  
plt.title('Recency by Cluster')  
plt.ylabel('Recency')  
  
plt.tight_layout()  
plt.show()
```



```
[58]: overlap_indices = monetary_outliers_df.index.intersection(frequency_outliers_df.
    ↪index)

monetary_only_outliers = monetary_outliers_df.drop(overlap_indices)
frequency_only_outliers = frequency_outliers_df.drop(overlap_indices)
monetary_and_frequency_outliers = monetary_outliers_df.loc[overlap_indices]

monetary_only_outliers["Cluster"] = -1
frequency_only_outliers["Cluster"] = -2
monetary_and_frequency_outliers["Cluster"] = -3

outlier_clusters_df = pd.concat([monetary_only_outliers, ↪
    ↪frequency_only_outliers, monetary_and_frequency_outliers])

outlier_clusters_df
```

```
[58]:      Customer ID      MonetaryValue  Frequency  LastInvoiceDate \
9          12357.00        11229.99       1 2010-11-16 10:05:00
25         12380.00        4782.84        4 2010-08-31 14:54:00
42         12409.00        12346.62       4 2010-10-15 10:24:00
48         12415.00        19468.84       4 2010-11-29 15:07:00
61         12431.00        4145.52       11 2010-12-01 10:03:00
...
4235        18223.00        7516.31       12 2010-11-17 12:20:00
4236        18225.00        7545.14       15 2010-12-09 15:46:00
4237        18226.00        6650.83       15 2010-11-26 15:51:00
4241        18231.00        4791.80       23 2010-10-29 14:17:00
4262        18260.00        7318.91       17 2010-11-30 12:25:00
```

	Recency	Cluster
9	23	-1
25	100	-1
42	55	-1
48	10	-1
61	8	-1
...	...	...
4235	22	-3
4236	0	-3
4237	13	-3
4241	41	-3
4262	9	-3

[476 rows x 6 columns]

```
[59]: cluster_colors = {-1: '#9467bd',
                      -2: '#8c564b',
                      -3: '#e377c2'}


plt.figure(figsize=(12, 18))

plt.subplot(3, 1, 1)
sns.violinplot(x=outlier_clusters_df['Cluster'], □
                ↪y=outlier_clusters_df['MonetaryValue'], palette=cluster_colors, □
                ↪hue=outlier_clusters_df["Cluster"])
sns.violinplot(y=outlier_clusters_df['MonetaryValue'], color='gray', □
                ↪linewidth=1.0)
plt.title('Monetary Value by Cluster')
plt.ylabel('Monetary Value')

plt.subplot(3, 1, 2)
sns.violinplot(x=outlier_clusters_df['Cluster'], □
                ↪y=outlier_clusters_df['Frequency'], palette=cluster_colors, □
                ↪hue=outlier_clusters_df["Cluster"])
sns.violinplot(y=outlier_clusters_df['Frequency'], color='gray', linewidth=1.0)
plt.title('Frequency by Cluster')
plt.ylabel('Frequency')

plt.subplot(3, 1, 3)
sns.violinplot(x=outlier_clusters_df['Cluster'], □
                ↪y=outlier_clusters_df['Recency'], palette=cluster_colors, □
                ↪hue=outlier_clusters_df["Cluster"])
sns.violinplot(y=outlier_clusters_df['Recency'], color='gray', linewidth=1.0)
plt.title('Recency by Cluster')
plt.ylabel('Recency')

plt.tight_layout()
plt.show()
```



```
[61]: cluster_labels = {
    0: "RETAIN",
    1: "RE-ENGAGE",
    2: "NURTURE",
    3: "REWARD",
    -1: "PAMPER",
    -2: "UPSELL",
    -3: "DELIGHT"
}
```

```
[62]: full_clustering_df = pd.concat([non_outliers_df, outlier_clusters_df])
full_clustering_df
```

```
[62]:      Customer ID      MonetaryValue  Frequency  LastInvoiceDate \
0          12346.00        169.36         2 2010-06-28 13:53:00
1          12347.00        1323.32        2 2010-12-07 14:57:00
2          12348.00        221.16         1 2010-09-27 14:59:00
3          12349.00        2221.14        2 2010-10-28 08:23:00
4          12351.00        300.93         1 2010-11-29 15:23:00
...
       ...
4235        18223.00        7516.31        12 2010-11-17 12:20:00
4236        18225.00        7545.14        15 2010-12-09 15:46:00
4237        18226.00        6650.83        15 2010-11-26 15:51:00
4241        18231.00        4791.80        23 2010-10-29 14:17:00
4262        18260.00        7318.91        17 2010-11-30 12:25:00

      Recency  Cluster
0        164      1
1         2      0
2         73      2
3         42      0
4         10      2
...
...
4235       22     -3
4236       0     -3
4237       13     -3
4241       41     -3
4262       9     -3

[4285 rows x 6 columns]
```

```
[63]: full_clustering_df["ClusterLabel"] = full_clustering_df["Cluster"].
    ↪map(cluster_labels)
full_clustering_df
```

```
[63]:          Customer ID      MonetaryValue  Frequency   LastInvoiceDate \
0            12346.00        169.36           2 2010-06-28 13:53:00
1            12347.00        1323.32          2 2010-12-07 14:57:00
2            12348.00        221.16           1 2010-09-27 14:59:00
3            12349.00        2221.14          2 2010-10-28 08:23:00
4            12351.00        300.93           1 2010-11-29 15:23:00
...
...          ...
4235         18223.00        7516.31          12 2010-11-17 12:20:00
4236         18225.00        7545.14          15 2010-12-09 15:46:00
4237         18226.00        6650.83          15 2010-11-26 15:51:00
4241         18231.00        4791.80          23 2010-10-29 14:17:00
4262         18260.00        7318.91          17 2010-11-30 12:25:00

      Recency  Cluster ClusterLabel
0        164       1    RE-ENGAGE
1         2       0     RETAIN
2        73       2    NURTURE
3        42       0     RETAIN
4        10       2    NURTURE
...
...        ...
4235       22      -3    DELIGHT
4236       0      -3    DELIGHT
4237       13      -3    DELIGHT
4241       41      -3    DELIGHT
4262       9      -3    DELIGHT

[4285 rows x 7 columns]
```

## 6 VISUALISATION

```
[64]: cluster_counts = full_clustering_df['ClusterLabel'].value_counts()
full_clustering_df["MonetaryValue per 100 pounds"] = ↴
    ↴full_clustering_df["MonetaryValue"] / 100.00
feature_means = full_clustering_df.groupby('ClusterLabel')[['Recency', ↴
    ↴'Frequency', 'MonetaryValue per 100 pounds']].mean()

fig, ax1 = plt.subplots(figsize=(12, 8))

sns.barplot(x=cluster_counts.index, y=cluster_counts.values, ax=ax1, ↴
    ↴palette='viridis', hue=cluster_counts.index)
ax1.set_ylabel('Number of Customers', color='b')
ax1.set_title('Cluster Distribution with Average Feature Values')


ax2 = ax1.twinx()

sns.lineplot(data=feature_means, ax=ax2, palette='Set2', marker='o')
```

```

ax2.set_ylabel('Average Value', color='g')
plt.show()

```



[ ]: